Probability and Statistics

Omid Safarzadeh

February 9, 2022

Table of contents

Moments definition

Definition 1.1

For each of integer n, the n^{th} moment of X is

$$\mu'_n = E[X^n].$$

The n^{th} central moment of X, μ_n, is

$$\mu_n = E[(X - \mu)^n],$$

where $\mu = \mu'_1 = E[X]$.

Expected Value

- Recall that"average" is an arithmetic average where all available observations are weighted equally.
- The expected value, on the other hand, is the average of all possible values a random variable can take, weighted by the probability distribution.
- The question is,which value would we expect the random variable to take on, on average.

Expected Value

Definition 1.2

The expected value or mean of a random variable $g(X)$, denoted by $E[g(X)]$, is

$$E[g(X)] = \begin{cases} \int_{-\infty}^{\infty} g(x) f_X(x) dx & \text{if } X \text{ is continuous} \\ \\ \sum_{x \in \mathcal{X}} g(x) f_X(x) = \sum_{x \in \mathcal{X}} g(x) P(X = x) & \text{if } X \text{ is discrete} \end{cases}$$

If $E[g(X)] = \infty$, we say that $E[g(X)]$ does not exist.

- we are taking the average of $g(x)$ over all of its possible values ($x \in \mathcal{X}$), where these values are weighted by the respective value of the pdf, $f_X(x)$.

Expected Value

Example 1.1

Suppose X has an exponential(λ) distribution, that is, it has pdf given by

$$f_X(x) = \frac{1}{\lambda}e^{-x/\lambda}, \quad 0 \leq x < \infty \quad \lambda > 0.$$

Then,

$$E[X] = \int_0^\infty \frac{1}{\lambda}xe^{-x/\lambda}dx = -xe^{-x/\lambda}|_0^\infty + \int_0^\infty e^{-x/\lambda}dx \quad (1)$$

$$= \int_0^\infty e^{-x/\lambda}dx = \lambda. \quad (2)$$

- To obtain this result, we use a method called integration by parts. This is based on
$$\int udv = uv - \int vdu.$$

Expected Value

- A very useful property of the expectation operator is that it is a linear operator.
- take a and b constants:

$$E[a + Xb] = a + E[Xb] = a + bE[x] = a + b\mu.$$

Expected Value

Theorem 1.1

Let X be a random variable and let a, b and c be constants. Then for any functions $g_1(x)$ and $g_2(x)$ whose expectations exist,

- $E[ag_1(X) + bg_2(X) + c] = aE[g_1(X)] + bE[g_2(X)] + c$.
- If $g_1(x) \geq 0$ for all x, then $E[g_1(X)] \geq 0$.
- If $g_1(x) \geq g_2(x)$ for all x, then $E[g_1(X)] \geq E[g_2(X)]$.
- If $a \leq g_1(x) \leq b$ for all x, then $a \leq E[g(X)] \leq b$.

Proof: Exercise!

Expected Value

Example 1.2

Let X have a uniform distribution, such that

$$f_X(x) = \begin{cases} 1 & \text{if } 0 \leq x \leq 1 \\ 0 & \text{if otherwise} \end{cases}$$

Define $g(X) = -\log X$. Then,

$$E[g(X)] = E[-\log X] = \int_0^1 -\log x \, dx = (-x \log x + x)|_0^1 = 1,$$

where we use integration by parts.

Variance

- variance measures the variation/dispersion/spread of the random variable around expectation.
- While the expectation is usually denoted by μ, σ^2 is generally used for variance.
- Variance is a second-order moment.

Variance

Definition 1.3

The variance of a random variable X is its second central moment,

$$Var(X) = E[(X - \mu)^2],$$

while $\sqrt{Var(X)}$ is known as the standard deviation of X.

- Importantly,
$$Var(X) = E[(X - \mu)^2] = E[X^2] - \mu^2.$$

covariance

- When it exists, the covariance of two random variables X and Y is defined as

$$Cov(X, Y) = E(\{X - E[X]\}\{Y - E[Y]\}).$$

Covariance and Correlation

- Let X and Y be two random variables. To keep notation concise, we will use the following notation.

$$E[X] = \mu_X, \quad E[Y] = \mu_Y, \quad Var(X) = \sigma_X^2 \quad and \quad Var(Y) = \sigma_Y^2.$$

Definition 2.1

The covariance of X and Y is the number defined by

$$Cov(X, Y) = E[(X - \mu_X)(Y - \mu_Y)].$$

Definition 2.2

The correlation of X and Y is the number defined by

$$\rho_{XY} = \frac{Cov(X, Y)}{\sigma_x \sigma_y},$$

which is also called the correlation coefficient.

Covariance and Correlation

- If large(small) values of X, tend to be observed with large(small) values of Y, then $Cov(X, Y)$ will be positive.

- Why so? Within the above setting, when $X > \mu_X$ then $Y > \mu_Y$ is likely to be true whereas when $X < \mu_X$ then $Y < \mu_Y$ is likely to be true. Hence

$$E[(X - \mu_X)(Y - \mu_Y)] > 0.$$

- Similarly, if large(small) values of X tend to be observed with small(large) values of Y, then will be negative.

Covariance and Correlation

- Correlation normalises covariance by the standard deviations and is, therefore, a more informative measure.
- If Cov(X, Y)=50 while Cov(W, Z)=0.9, this does not necessarily mean that there is a much stringer relationship between X and Y. For example, if Car(X)=Var(Y)=100 while Var(W)=Var(Z)=1, then

$$\rho_{XY} = 0.5 \quad \rho_{WZ} = 0.9.$$

Covariance

Theorem 2.1

For any random variables X and Y,

$$Cov(X, Y) = E[XY] - \mu_X \mu_Y.$$

- **Proof**: Exercise!

Covariance and Correlation

Theorem 2.2

If $X \perp\!\!\!\perp Y$, then Cov$(X, Y) = \rho_{XY} = 0$.

- **Proof**: Exercise!
- It is crucial to note that although $X \perp\!\!\!\perp Y$ implies that $Cov(X, Y) = \rho_{XY} = 0$, the relationship does not necessarily hold in the reverse direction.

Covariance and Correlation

Theorem 2.3

If X and Y are any two random variables and a and b are any two constants, then

$$Var(aX + bY) = a^2 Var(X) + b^2 Var(Y) + 2ab Cov(X, Y)$$

If X and Y are independent random variables, then

$$Var(aX + bY) = a^2 Var(X) + b^2 Var(Y).$$

- **Proof**: Exercise!

Covariance and Correlation

- Note that if two random variables, X and Y, are positively correlated, then

$$Var(X + Y) > Var(X) + Var(Y),$$

  whereas if X and Y are negatively correlated, then

$$Var(X + Y) < Var(X) + Var(Y).$$

- For positively correlated random variables, large values in one tend to be accompanied by large values in the other. Therefore, the total variance is magnified.

- Similarly, for negatively correlated random variables, large values in one tend to be accompanied by small values in the other. Hence, the variance of the sum is dampened.

Variance of Sums of Random Variables

- Let a_i be some constant and X_i be some random variable, where $i = 1, ..., n$.
- Then

$$Var(\sum_{i=1}^{n} a_i X_i) = \sum_{i=1}^{n} a_i^2 Var(X_i) + \sum_{i \neq j} \sum a_i a_j Cov(X_i, X_j).$$

third and fourth moments

- third and fourth moments are concerned with how symmetric and fat-tailed the underlying distribution is.

- moment generating function can be used to obtain moments of a random variable.

Moments and Moment Generating Functions

Definition 3.1

Let X be a random variable with cdf F_X. The moment generating function (mgf) of X (or F_X), denoted by $M_X(t)$, is

$$M_X(t) = E[e^{tX}],$$

provided that the expectation exists for t in some neighbourhood of 0. That is, there is an $h > 0$ such that, for all t in $-h < t < h$, $E[e^{tX}]$ exists. If the expectation does not exist in a neighbourhood of 0, we say that the mgf does not exist.

- We can write the mgf of X as

$$M_X(t) = \int_{-\infty}^{\infty} e^{tx} f_X(x) dx \quad \text{if } X \text{ is continuous,}$$

$$M_X(t) = \sum_x e^{tx} P(X = x) \quad \text{if } X \text{ is discrete.}$$

Moment Generating Functions

Theorem 3.1

If X has mgf $M_X(t)$, then

$$E[X^n] = M_X^{(n)}(0),$$

where we define

$$M_X^{(n)}(0) = \frac{d^n}{dt^n} M_X(t)|_{t=0}.$$

That is, the n^{th} moment is equal to the n^{th} derivative of $M_X(t)$ evaluated at t=0.

Normal mgf

- Now consider the pdf for $X \sim N(\mu, \sigma^2)$

$$f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} exp[-\frac{1}{2}(\frac{x-\mu}{\sigma})^2], \quad -\infty < x < \infty.$$

- The mgf is given by

$$M_X(t) = E[e^{Xt}] = \int_{-\infty}^{\infty} e^{tx} \frac{1}{\sqrt{2\pi\sigma^2}} e^{-(x-\mu)^2/2\sigma^2} dx.$$

Normal mgf

- Note that:

$$M_X(t) = exp(\mu t + \frac{\sigma^2 t^2}{2}).$$

- Proof: Exercise!
- Clearly,

$$E[X] = \frac{d}{dt} M_X(t)|_{t=0} = (\mu + \sigma^2 t) exp(\mu t + \frac{\sigma^2 t^2}{2})|_{t=0} = \mu,$$

$$E[X^2] = \frac{d^2}{dt^2} M_X(t)|_{t=0} = \sigma^2 exp(\mu t + \frac{\sigma^2 t^2}{2}|_{t=0}$$

$$+ (\mu + \sigma^2 t)^2 exp(\mu t + \frac{\sigma^2 t^2}{2})^2|_{t=0}$$

$$= \sigma^2 + \mu^2,$$

$$Var(X) = E[X^2] - \{E[X]\}^2 = \sigma^2 + \mu^2 - \mu^2 = \sigma^2.$$

Matrix Notation for Moments

- Now, let X and Y be $(r*1)$ and $(c*1)$ random vectors, respectively. Define
- In other words,

$$Cov(X, Y) = \begin{bmatrix} Cov(X_1, Y_1) & \cdots & Cov(X_1, Y_c) \\ \vdots & \ddots & \vdots \\ Cov(X_r, Y_1) & \cdots & Cov(X_r, Y_c) \end{bmatrix}$$

$$= E \begin{bmatrix} \{X_1 - E[X_1]\}\{Y_1 - E[Y_1]\} & \cdots & \{X_1 - E[X_1]\}\{Y_c - E[Y_c]\} \\ \vdots & \ddots & \vdots \\ \{X_r - E[X_r]\}\{Y_1 - E[Y_1]\} & \cdots & \{X_r - E[X_r]\}\{Y_c - E[Y_c]\} \end{bmatrix}$$

Matrix Notation for Moments

$$= E \left[\begin{pmatrix} X_1 - E[X_1] \\ \vdots \\ X_r - E[X_r] \end{pmatrix} (Y_1 - E[Y_1], \cdots, Y_c - E[Y_c]) \right],$$

$$= E(\{X - E[X]\}\{Y - E[Y]\}').$$

Matrix Notation for Moments

- Usually, for a $(c * 1)$ vector X, one would write $Cov(X)$ for $Cov(X, X)$,
- This is given by

$$= Cov(X) \begin{bmatrix} Var(X_1) & \cdots & Cov(X_1, X_c) \\ \vdots & \ddots & \vdots \\ Cov(X_1, X_c) & \cdots & Var(X_c) \end{bmatrix},$$

which is a $(c * c)$ symmetric matrix.

Matrix Notation for Moments

- We can also consider block structures. Let

$$X = \begin{pmatrix} Y \\ Z \end{pmatrix},$$

where Y is $(p * 1)$ vector and Z is a $(q * 1)$ vector.

- Then,

$$Cov(X) = E(\{\begin{pmatrix} Y \\ Z \end{pmatrix} - E[\begin{pmatrix} Y \\ Z \end{pmatrix}]\}\{\begin{pmatrix} Y \\ Z \end{pmatrix} - E[\begin{pmatrix} Y \\ Z \end{pmatrix}]\}')$$

$$= E \begin{pmatrix} \{Y - E[Y]\}\{Y - E[Y]\}' & \{Y - E[Y]\}\{Z - E[Z]\}' \\ \{Z - E[Z]\}\{Y - E[Y]\}' & \{Z - E[Z]\}\{Z - E[Z]\}' \end{pmatrix}$$

$$= \begin{pmatrix} Cov(Y) & Cov(Y, Z) \\ Cov(Z, Y) & Cov(Z) \end{pmatrix},$$

where $Cov(Y)$ is $(p * p)$, $Cov(Y, Z)$ is $(p * q)$, $Cov(Z, Y)$ is $(q * p)$ and $Cov(Z)$ is $(q * q)$.

Matrix Notation for Moments

- Let a and b be $(r*1)$ and $(c*1)$ non-stochastic vectors. We might encounter terms such as $Cov(a'X, b'Y)$ or $Var(a'X)$.
- Let $E[X_i] = \mu_{X_i}$, $E[Y_i] = \mu_{Y_i}$ and $Cov(X_i, Y_j) = \Sigma_{X_i, Y_j}$. Then

$$Cov(a'X, b'Y) = Cov(\sum_{i=1}^{r} a_i X_i, \sum_{j=1}^{c} b_j Y_j)$$

$$= E\{[\sum_{i=1}^{r} a_i(X_i - \mu_{X_i})][\sum_{j=1}^{c} b_j(Y_j - \mu_{Y_j})]\}$$

$$= \sum_{i=1}^{r} \sum_{j=1}^{c} a_i b_j E[(X_i - \mu_{X_i})(Y_j - \mu_{Y_j})]$$

$$= \sum_{i=1}^{r} \sum_{j=1}^{c} a_i b_j \Sigma_{X_i, Y_j} = a'\Sigma_{XY} b = a' Cov(X, Y) b.$$

Matrix Notation for Moments

- Now, let $\Sigma_{ij} = Cov(X_i, X_j)$ and $\Sigma_{XX} = Var(X)$. Then,

$$Var(a'X) = E[(\sum_{i=1}^{r} a_i X_i - E[\sum_{i=1}^{r} a_i X_i])^2]$$

$$= E\{[\sum_{i=1}^{r} a_i(X_i - \mu_i)][\sum_{i=1}^{r} a_i(X_i - \mu_i)]\}$$

$$= \sum_{i=1}^{r} \sum_{j=1}^{r} a_i a_j E[(X_i - \mu_i)(X_j - \mu_j)]$$

$$= \sum_{i=1}^{r} \sum_{j=1}^{r} a_i a_j \Sigma_{ij} = a' Var(X)a.$$

Matrix Notation for Moments

- Now, Consider

$$Var(X + Y) = E\{[(X - \mu_X) + (Y - \mu_Y)][(X - \mu_X) + (Y - \mu_Y)]'\}$$

$$= E[(X - \mu_X)(X - \mu_X)'] + E[(X - \mu_X) + (Y - \mu_Y)]'$$

$$+ E[(Y - \mu_Y)(X - \mu_X)'] + E[(Y - \mu_Y) + (Y - \mu_Y)]'$$

$$= \Sigma_{XX} + \Sigma_{XY} + \Sigma_{YX} + \Sigma_{YY}.$$

- Using this, we get

$$Var[a'(X + Y)] = a'(\Sigma_{XX} + \Sigma_{XY} + \Sigma_{YX} + \Sigma_{YY})a$$

$$= a'\Sigma_{XX}a + 2a'\Sigma_{XY}a + a'\Sigma_{YY}a,$$

where we use the fact that

$$a'\Sigma_{XY}a = a'\Sigma_{YX}a$$

Matrix Notation for Moments

- These results easily extend to cases where a and b are replaced by matrices.

$$E[RX] = RE[X]$$

$$
\begin{aligned}
Var(RX) &= E[R(X - \mu_X)(X - \mu_X)'R'] \\
&= RE[(X - \mu_X)(X - \mu_X)']R' \\
&= R\Sigma_{XX}R'.
\end{aligned}
$$

Discrete Uniform Distribution

- A random variable X has a discrete *uniform*(1,N) distribution if

$$P(X = x|N) = \frac{1}{N}, \quad x = 1, 2, ..., N,$$

where N is a specified integer. This distribution puts equal mass on each of the outcomes $1, 2, ..., N$.

- $E(X)=(N+1)/2$
- $Var(X)=(N+1)(N-1)/2$

Bernoulli Distribution

- A random variable X has Bernoulli(p) distribution if

$$X = \begin{cases} 1 & \text{with probability } p \\ 0 & \text{with probability } 1-p \end{cases} \qquad 0 \le p \le 1.$$

- $X = 1$ is often termed as "success" and p is, accordingly, the probability of success. Similarly, $X = 0$ is termed a "failure".

- Now,

$$E[X] = 1 * p + 0 * (1 - p) = p,$$

$$\text{and} \quad Var(X) = (1 - p)^2 p + (0 - p)^2 (1 - p) = p(1 - p).$$

Binomial Distribution

- This is based on a Bernoulli trial which is an experiment with two, and only, two, possible outcomes.
- Assume, we have n trials of a Bernoulli distribution, and we are interested to probability of having y results as success. It means than n-y times we had failure. Also assume that these events are independent of each other. Hence: the distribution of the total number of successes in n trials is Binomial Distribution
- Examples:
  1. Tossing a coin ($p =$ probability of a head, $X = 1$ if heads)
  2. Election polls ($X = 1$ if candidate A gets vote)
  3. Probability of Default Risk ($p =$ probability that a person defaults in his loan payments)
  4. in ML we use it to construct Binary Cross-Entropy Loss Function

Binomial Distribution

- Take $Y =$ "total number of successes in n trials"
- There are many possible orderings of the events that would lead to this outcome. Any particular such ordering has probability

$$p^y(1-p)^{n-y}.$$

- Since there are $\binom{n}{y}$ such sequences, we have

$$P(Y = y|n, p) = \binom{n}{y} p^y (1-p)^{n-y}, \quad y = 0, 1, ..., n,$$

and Y is called a *binomial(n,p)* random variable.

- $E[X] = np$
- $Var(X) = np(1-p)$ (**Proof**: Exercise!)

Poisson Distribution

- In modelling a phenomenon in which we are waiting for an occurrence (such as waiting for a bus), the number of occurrence in a given time interval can be modelled by the Poisson distribution.
- The basic assumption is as follows: for small time intervals, the probability of an arrival is proportional to the length of waiting time.
- If we are waiting for the bus, the probability that a bus will arrive within the next hour is higher than the probability that it will arrive within 5 minutes.
- Other possible applications are distribution of bomb hits in an area or distribution of fish in a lake.
- The only parameter is λ, also sometimes called the "intensity parameter."

Poisson Distribution

- $P(X = x | \lambda) = \frac{e^{-\lambda} \lambda^x}{x!}, x = 0, 1, ...$
- $E[X] = \lambda$
- $Var(X) = \lambda$
- **Proof**: Exercise!

Poisson Distribution

Example 5.1

As an example of a waiting-for-occurrence application, consider a telephone operator who, on average, handles fire calls every 3 minutes. What is the probability that there will be no calls in the next minute? At least two calls? If we let $X =$number of calls in a minute, then X has a Poisson distribution with $E[X] = \lambda = 5/3$. So,

$$P(\text{no calls in the next minute}) = P(X = 0)$$

$$= \frac{e^{-5/3}(5/3)^0}{0!} = e^{-5/3} = 0.189$$

and $\qquad P(\text{at least two calls in the next minute}) = P(X \geq 2)$

$$= 1 - P(X = 0) - P(X = 1)$$

$$= 1 - 0.189 - \frac{e^{-5/3}(5/3)^1}{1!}$$

$$= 0.496.$$

Uniform Distribution

- The continuous uniform distribution is defined by spreading mass uniformly over an interval $[a, b]$. Its pdf is given by

$$f(x|a, b) = \begin{cases} \frac{1}{b-a} & \text{if } x \in [a, b] \\ 0 & \text{if otherwise} \end{cases}.$$

- One can easily show that

$$\int_a^b f(x)dx = 1,$$

$$E[X] = \frac{b+a}{2},$$

$$Var(X) = \frac{(b-a)^2}{12}.$$

- In many cases, when people say Uniform distribution, they implictly mean $(a, b) = (0, 1)$.

Exponential Distribution

- pdf of Exponential Distribution :

$$f(x|\beta) = \frac{1}{\beta} e^{-e/\beta}, \quad 0 < x < \infty.$$

- we have

$$E[X] = \beta \quad \text{and} \quad Var(X) = \beta^2$$

- this distribution is that it has no memory.

Exponential Distribution

- If $X \sim$ exponential(β), then, for $s > t \geq 0$,

$$P(X > s | X > t) = \frac{P(X > s, X > t)}{P(X > t)} = \frac{P(X > s)}{P(X > t)}$$

$$= \frac{\int_s^\infty \frac{1}{\beta} e^{-x/\beta} dx}{\int_t^\infty \frac{1}{\beta} e^{-x/\beta} dx} = \frac{e^{-s/\beta}}{e^{-t/\beta}}$$

$$= e^{-(s-t)/\beta} = P(X > s - t).$$

- This is because,

$$\int_{s-t}^\infty \frac{1}{\beta} e^{-x/\beta} dx = -e^{-x/\beta}|_{s-t}^\infty = e^{-(s-t)/\beta}.$$

- What does this mean? When calculating $P(X > s | X > t)$, what matters is not whether X has passed a threshold or not. What matters is the distance between the threshold and the value to be reached.
- If Mr X has been fired more than 10 times, what is the probability that he will be fired more than 12 times? It is not different from the probability that a person, who has been fired once, will be fired more than two times. History does not matter.

Normal Distribution

- We now consider the normal distribution or the Gaussian distribution.
- Why is this distribution so popular?
  1. Analytical tractability
  2. Bell shaped or symmetric
  3. It is central to Central Limit Theorem; this type of results guarantee that, under (mild) conditions, the normal distribution can be used to approximate a large variety of distribution in large samples.
- The pdf is given by,

$$f(x|\mu, \sigma^2) = \frac{1}{\sigma\sqrt{2\pi}} exp[-\frac{(x-\mu)^2}{2\sigma^2}].$$

Normal Distribution

- This distribution is usually denoted as $N(\mu, \sigma^2)$.
- A very useful result is that for $X \sim N(\mu, \sigma^2)$,

$$Z = \frac{X - \mu}{\sigma} \sim N(0, 1).$$

- $N(0, 1)$ is known as the standard normal distribution.
- To see this, consider the following:

$$P(Z \leq z) = P(\frac{(X - \mu)}{\sigma} \leq z)$$

$$= P(X \leq z\sigma + \mu)$$

$$= \frac{1}{\sigma\sqrt{2\pi}} \int_{-\infty}^{z\sigma+\mu} e^{-(x-\mu)^2/2\sigma^2} dx$$

$$= \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{z} e^{-t^2/2} dt,$$

where we substitute $t = (x - \mu)/\mu$. Notice that this implies that $dt/dx = 1/\sigma$. This shows that $P(Z \leq z)$ is the standard normal cdf.

Lognormal Distribution

- Let X be a random variable such that

$$\log X \sim N(\mu, \sigma^2).$$

Then, X is said to have a lognormal distribution.

- By using a transformation Theorem , the pdf of X is given by,

$$f(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \frac{1}{x} exp[-\frac{(\log x - \mu)^2}{2\sigma^2}],$$

where $0 < x < \infty$, $-\infty < \mu < \infty$, and $\sigma > 0$.

Laplace distribution

- If $X \sim Lap(\mu, b)$,

$$f(x|\mu, b) = \frac{1}{2b} exp(-\frac{|x - \mu|}{b})$$

- then $E[X] = \mu$, $Var(X) = 2b^2$
- The Lasso Regression is sort of a Bayesian regression with a Laplacian prior
- Laplace is applied to extreme events like rainfalls, river discharges

Reference

Casella, G., & Berger, R. (2002). *Statistical inference*. Cengage Learning. https://books.google.fr/books?id=FAUVEAAAQBAJ